

Bayesian data analysis – Assignment 3

General information

- The recommended tool in this course is R (with the IDE R-Studio). You can download R [here](#) and R-Studio [here](#). There are tons of tutorials, videos and introductions to R and R-Studio online. You can find some initial hints from [RStudio Education pages](#).
- When working with R, we recommend writing the report using R markdown and the provided [R markdown template](#). The template includes the formatting instructions and how to include code and figures.
- Instead of R markdown, you can use other software to make the pdf report, but the same instructions for formatting should be used. These instructions are available also in [the PDF produced from the R markdown template](#).
- Report all results in a single, **anonymous** *.pdf -file and return it to [peergrade.io](#).
- The course has its own R package `aaltobda` with data and functionality to simplify coding. To install the package just run the following (upgrade="never" skips question about updating other packages):
 1. `install.packages("remotes")`
 2. `remotes::install_github("avehtari/BDA_course_Aalto",
subdir = "rpackage", upgrade="never")`
- Many of the exercises can be checked automatically using the R package `markmyassignment`. Information on how to install and use the package can be found [here](#). There is no need to include `markmyassignment` results in the report.
- Recommended additional self study exercises for each chapter in BDA3 are listed in the course web page.
- We collect common questions regarding installation and technical problems in a course Frequently Asked Questions (FAQ). This can be found [here](#).
- Deadlines for all assignments can be found on the course web page and in [peergrade.io](#).
- If you have any suggestions or improvements to the course material, please post in the course chat feedback channel, create an issue, or submit a pull request to the public repository!

Information on this assignment

This assignment is related to Chapters 2 and 3. The maximum amount of points from this assignment is 9. Use [Frank Harrell's recommendations](#) on how to state results in Bayesian two group comparisons (and note that there is no point null hypothesis testing in this assignment).

Reading instructions: Chapter 2 and 3 in BDA3, see reading instructions [here](#) and [here](#).

Grading instructions: The grading will be done in peergrade. All grading questions and evaluations for assignment 3 can be found [here](#).

To use markmyassignment for this assignment, run the following code in R:

```
> library(markmyassignment)
> assignment_path <-
  paste("https://github.com/avehtari/BDA_course_Aalto/",
        "blob/master/assignments/tests/assignment3.yml", sep="")
> set_assignment(assignment_path)
> # To check your code/functions, just run
> mark_my_assignment()
```

1. Inference for normal mean and deviation (3 points)

A factory has a production line for manufacturing car windshields. A sample of windshields has been taken for testing hardness. The observed hardness values \mathbf{y}_1 can be found in file `windshields1.txt`. The data can also be accessed from the `aaltobda` R package as follows:

```
> library(aaltobda)
> data("windshields1")
> head(windshields1)

[1] 13.357 14.928 14.896 15.297 14.820 12.067
```

We may assume that the observations follow a normal distribution with an unknown standard deviation σ . We wish to obtain information about the unknown average hardness μ . For simplicity we assume standard uninformative prior discussed in the book, that is, $p(\mu, \sigma^2) \propto (\sigma^2)^{-1}$. It is not necessary to derive the posterior distribution in the report, as it has already been done in the book.

Below are test examples that can be used. The functions below can also be tested with `markmyassignment`. **Note!** This is *only* a test case. You need to change to the full data `windshields` above when reporting your results.

```
> windshields_test <- c(13.357, 14.928, 14.896, 14.820)
```

In the report, formulate (1) model likelihood, (2) the prior, and (3) the resulting posterior.

- a) What can you say about the unknown μ ? Summarize your results using Bayesian point estimate (i.e. $E(\mu|y)$), a posterior interval (95%), and plot the density. A test example can be found below for an uninformative prior. **Note!** Posterior intervals are also called credible intervals and are different from confidence intervals.

```
> mu_point_est(data = windshields_test)

[1] 14.5

> mu_interval(data = windshields_test, prob = 0.95)

[1] 13.3 15.7
```

- b) What can you say about the hardness of the next windshield coming from the production line before actually measuring the hardness? Summarize your results using Bayesian point estimate, a *predictive* interval (95%), and plot the density. A test example can be found below.

```
> mu_pred_point_est(data = windshields_test)

[1] 14.5

> mu_pred_interval(data = windshields_test, prob = 0.95)

[1] 11.8 17.2
```

Note! Predictive intervals are different from posterior intervals.

Hint With a conjugate prior a closed form posterior is Student's t form (see equations in the book). R users can use the `dt` function after doing input normalisation. We have added an R function `dtnew()` in the `aaltobda` R package which does that. For generating samples, you can use the corresponding `rtnew` function.

2. Inference for the difference between proportions (3 points)

An experiment was performed to estimate the effect of beta-blockers on mortality of cardiac patients. A group of patients was randomly assigned to treatment and control groups: out of 674 patients receiving the control, 39 died, and out of 680 receiving the treatment, 22 died. Assume that the outcomes are independent and binomially distributed, with probabilities of death of p_0 and p_1 under the control and treatment, respectively. Set up a noninformative or weakly informative prior distribution on (p_0, p_1) .

In the report, formulate (1) model likelihood, (2) the prior, and (3) the resulting posterior.

- a) Summarize the posterior distribution for the odds ratio, $(p_1/(1 - p_1))/(p_0/(1 - p_0))$. Compute the point estimate, a posterior interval (95%), and plot the histogram. Use [Frank Harrell's recommendations](#) how to state results in Bayesian two group comparison. Below is a test case on how the odd ratio should be computed. **Note!** This is *only* a test case. You need to change to the real posteriors when reporting your results.

```
> set.seed(4711)
> p0 <- rbeta(100000, 5, 95)
> p1 <- rbeta(100000, 10, 90)
> posterior_odds_ratio_point_est(p0 = p0, p1 = p1)
[1] 2.676

> posterior_odds_ratio_interval(p0 = p0, p1 = p1, prob = 0.9)
[1] 0.875 6.059
```

- b) Discuss the sensitivity of your inference to your choice of prior density with a couple of sentences.

Hint With a conjugate prior, a closed-form posterior is the Beta form for each group separately (see equations in the book). You can use `rbeta()` to sample from the posterior distributions of p_0 and p_1 , and use these samples and odds ratio equation to get samples from the distribution of the odds ratio.

3. Inference for the difference between normal means (3 points)

Consider a case where the same factory has two production lines for manufacturing car windshields. Independent samples from the two production lines were tested for hardness. The hardness measurements for the two samples \mathbf{y}_1 and \mathbf{y}_2 are given in the files `windshields1.txt` and `windshields2.txt`. These can be accessed directly with

```
> data("windshields1")
> data("windshields2")
```

We assume that the samples have unknown standard deviations σ_1 and σ_2 .

In the report, formulate (1) model likelihood, (2) the prior, and (3) the resulting posterior.

Use uninformative or weakly informative priors and answer the following questions:

- a) What can you say about $\mu_d = \mu_1 - \mu_2$? Summarize your results using a Bayesian point estimate, a posterior interval (95%), and plot the histogram. Use [Frank Harrell's recommendations](#) how to state results in Bayesian two group comparison.
- b) What is the probability that the means are the same? Explain your reasoning with a couple of sentences.

Hint With a conjugate prior, a closed-form posterior is Student's t form for each group separately (see equations in the book). You can use `rt()` function to sample from the posterior distributions of μ_1 and μ_2 , and use these samples to get samples from the distribution of the difference $\mu_d = \mu_1 - \mu_2$. The equivalent function in R is the `rt` function. Be careful to scale them and shift them according to their mean and variance values in R, as described above.

Hint Posterior distributions of μ_1 and μ_2 are continuous, and thus the posterior distribution of the difference $\mu_d = \mu_1 - \mu_2$ is also continuous. What is the probability that $\mu_d = 0$?